

# Interprocess Communication (Part 1)

Operating Systems

Assoc. Prof. Milos Jovanovik, PhD

# Agenda

- ▶ Problems in IPC
  - Race Condition
  - Critical Section
  - Deadlocks
- ▶ Solution Techniques
  - Lock variables
  - Hardware solutions
  - Semaphores
  - Monitors
- ▶ Classical IPC Problems



# Interprocess Communication

- ▶ Sometimes, processes need to communicate
  - How can they do this?
- ▶ How can they not get in each others way when sharing resources?
  - Two clients trying to reserve the last airplane seat;
- ▶ How can they sequence their actions when there are dependencies?
  - Process A needs to provide something first, so that Process B can read and process it afterwards;

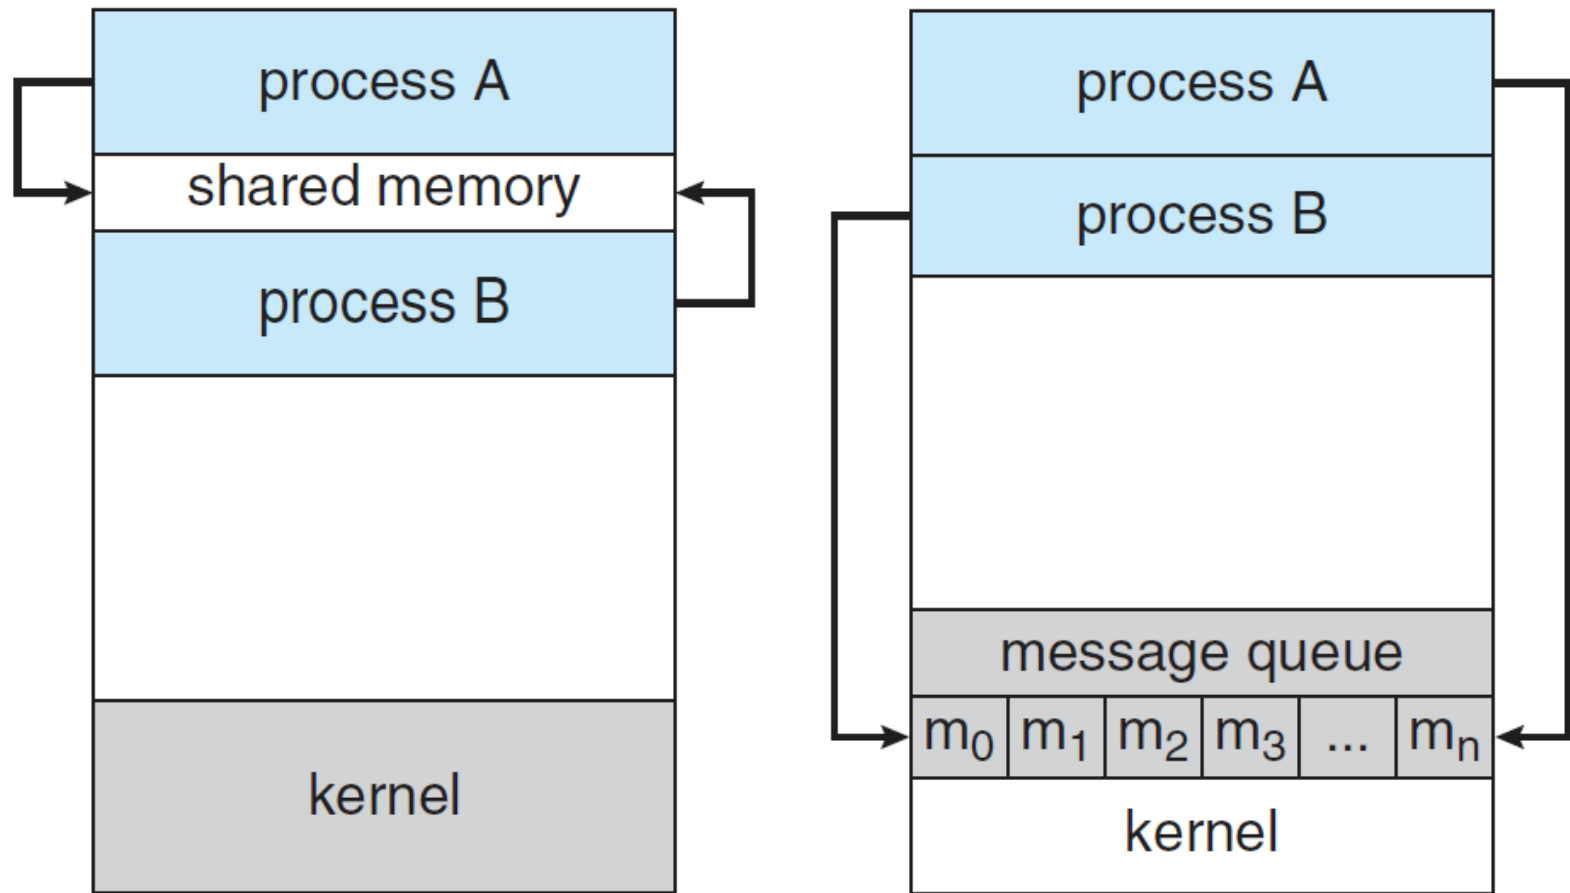


# Interprocess Communication

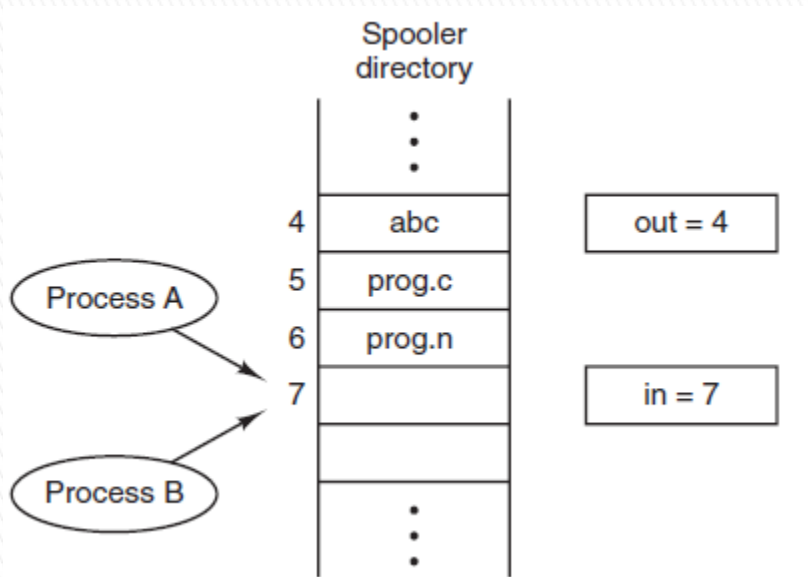
- ▶ Processes executing concurrently in the operating system may be:
  - **Independent**: they cannot affect or be affected by other processes;
  - **Cooperating**: they can affect or be affected by other processes;
    - Any process sharing data with other processes is a cooperating process
- ▶ Models of interprocess communication:
  - Shared memory
  - Message passing



# Interprocess Communication



# IPC using a Shared Variable



- ▶ Two shared variables:
  - **out** – next file to be printed
  - **in** – next free slot
- ▶ Process A and B want to print
- ▶ Process A reads the variable **in** and memorizes it in local variable **next-free**
- ▶ A clock interrupt occurs and the CPU decides to switch to process B
- ▶ Process B reads the value 7 from **in**, sends the file name and it sets **in** to 8
- ▶ Process A activates again and sends the file name to the place pointed by the **next-free**, in location 7, “overriding” the file of process B
- ▶ User B will wait forever for the printed document

**Problem: Process B starts to use the shared variable before A finishes with the variable!**

# Race Condition

- ▶ Situations where two or more processes are reading or writing shared data and the final result depends on who runs precisely when, is called a **race condition** between the processes
- ▶ Race conditions are difficult to detect
  - Most test runs may be fine, with a race condition happening only once in a while;
- ▶ The user needs an OS mechanism to avoid race conditions – main design goal of OS
- ▶ Solution: mutual exclusion
  - Making sure that if one process is using a shared variable or file, the other processes will be excluded from doing the same thing.



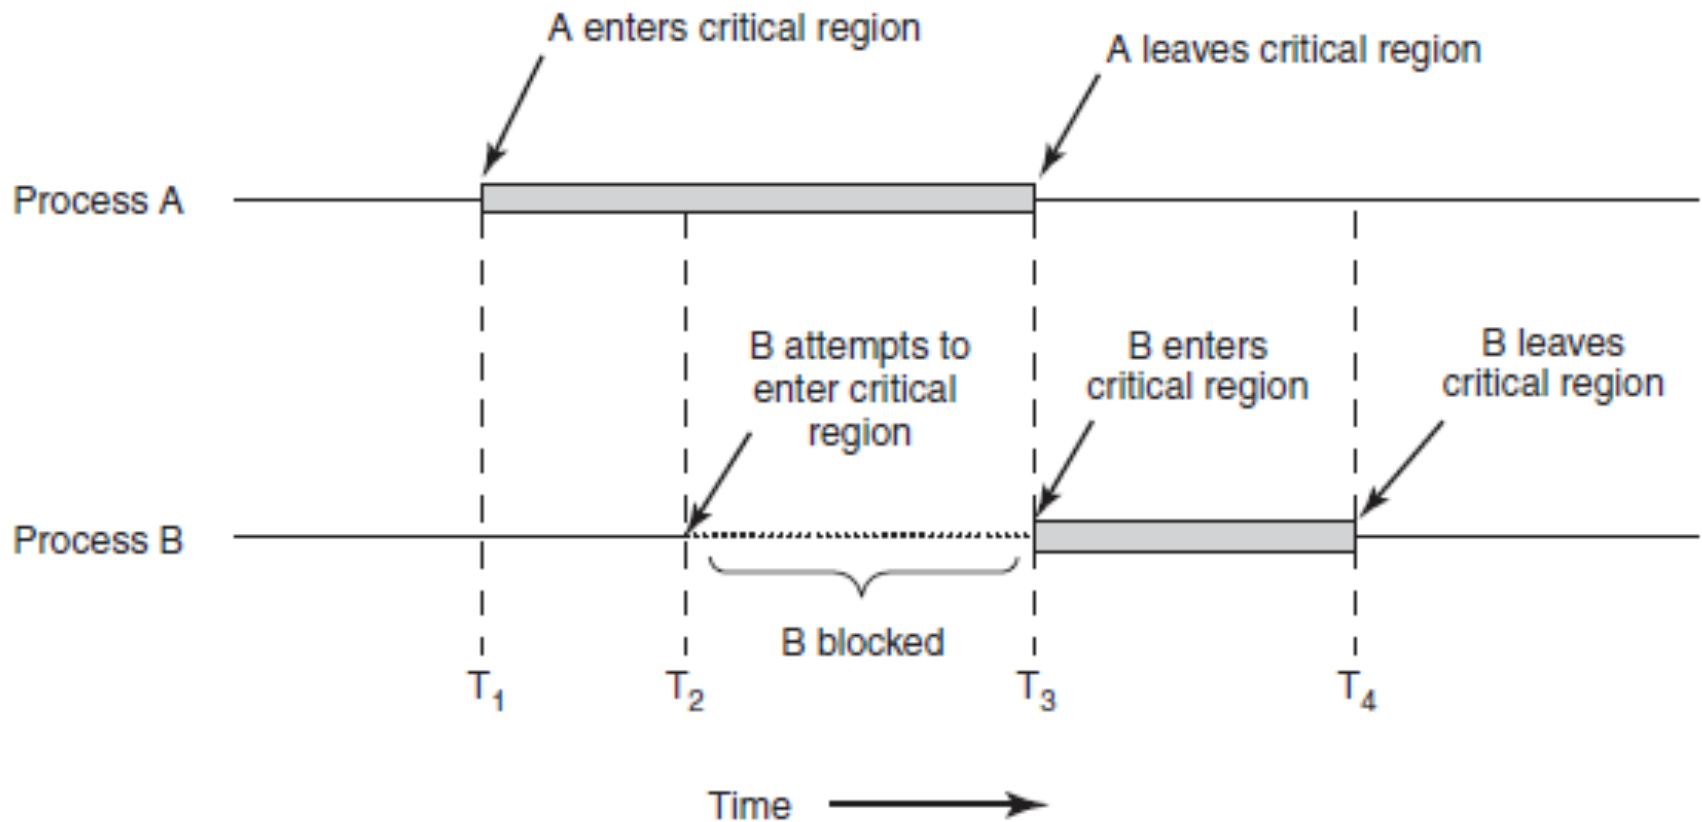
# Critical Region / Segment

- ▶ A critical region is a part of the program where the shared memory (variables), files, etc., are being accessed
- ▶ It's an abstraction which consists of a number of sequential program statements, which have to be executed without interrupt and errors
  - No two processes should be in their critical section at the same time, so that they avoid race conditions





# Critical Sections and Mutual Exclusion



# Critical Sections


shared double balance;

Code for P1

```
    . . .  
balance += amount;  
    . . .
```

Code for P2

```
    . . .  
balance -= amount;  
    . . .
```



# Critical Sections

shared double balance;

Code for P1

```
    . . .  
balance += amount;  
    . . .
```

Process P1 (assembler)

```
load R1, balance  
load R2, amount  
add R1, R2  
store R1, balance
```

Code for P2

```
    . . .  
balance -= amount;  
    . . .
```

Process P2 (assembler)

```
load R1, balance  
load R2, amount  
sub R1, R2  
store R1, balance
```



# Race Conditions

- ▶ There is a race between the processes – who will execute its critical section
- ▶ The section can be defined in different code structures, in different processes
  - A static analytics will not detected possible race conditions
- ▶ Without mutual exclusion, multiple executions of the same programs may result with **inconsistent results**
- ▶ There is a need for an OS mechanism for the user, to be able to handle race conditions



# Critical Section Problem

- ▶ A system with  $n$  processes  $\{P_0, P_1, \dots, P_{n-1}\}$
- ▶ Each process has a critical section (it changes variables' values, tables, writes to files shared with other processes, ...)
- ▶ RULE:
  - When one process works in its critical section, no other process has permission to execute code in its own critical section;
- ▶ We need a protocol for process cooperation



# General Structure of $P_i$

- ▶ Each process has to ask for a permission in order to enter in its critical section
- ▶ The critical section can be followed with an exit section

```
do {  
    entry_section  
    critical_section  
    exit_section  
    noncritical_section  
}
```



# Possible OS mechanisms

- ▶ Hardware solutions
  - Disabling interrupts
  - Special atomic instructions
- ▶ Software solutions
  - Lock variables
  - Semaphores
  - Monitors
- ▶ ...



# Implementation Criteria

- ▶ Each implementation needs to be:
  - **Correct:** only one process may execute the critical section at any given moment.
  - **Efficient:** entering and exiting a critical section has to be fast.
  - **Flexible:** a good implementation allows a maximized concurrency and has minimal limitations.





# Correct Solution

- ▶ General conditions for a correct solution of a racing condition:
  1. Mutual exclusion: No two processes may be simultaneously inside their critical regions.
  2. Progress: If no process is in the critical section, and there are processes waiting to enter it, one of them must get access.
  3. Limited waiting: No process should have to wait forever to enter its critical region.



# 1. Shared Lock Variable – Busy Waiting

A  
l  
g  
o  
r  
i  
t  
h  
m  
1

- ▶ Processes ( $P_0$  and  $P_1$ ) are sharing the variable **turn**
- ▶ A process can access a variable when its turn comes
- ▶ When a process wants to access a variable, but it is not its turn, it is in a state called **busy waiting** (wasting CPU time)

```
while (TRUE) {  
    while (turn != 0)    /* loop */ ;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)    /* loop */ ;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)



# Advantages and Problems

- ▶ This solution allows only one process to be in the critical section in a given moment.
- ▶ It does not allow progress: we need strict process alternation of the processes which are executing the critical section.
- ▶ **A problem arises if one process is much slower than the other.**

Process 0	Process 1
	turn = 0
wait for turn = 0	noncritical_region (enter)
critical_region (enter)	
critical_region (finish)	
turn = 1	
noncritical_region (enter)	
noncritical_region (finish)	
wait for turn = 0 (rule 2 is broken!)	

# Algorithm 2

- ▶ The first algorithm does not keep all the necessary information for the process state, it keeps only information about which process can enter in its critical section
- ▶ We introduce a Boolean array  $\text{flag}[P_1..P_2]$
- ▶ If  $\text{flag}[P_i]$  is true, the  $i^{\text{th}}$  process  $P_i$  is ready to enter its critical section
- ▶ The  $i^{\text{th}}$  process checks if the  $j^{\text{th}}$  process is NOT ready to enter its critical section
- ▶ If also  $\text{flag}[P_j]$  is true, then  $P_i$  waits until  $\text{flag}[P_j]$  becomes false, then  $P_i$  will enter its critical section



# Solution: Locking

Flag[P1..P2]

Process 1

```
Flag[P1] = true;
while(Flag[P2])
    ; // wait
/* enter C.S. */
/* leave C.S. */
Flag[P1] = false;
```

.....

Process 2

```
Flag[P2] = true;
while(Flag[P1])
    ; // wait
/* enter C.S. */
/* leave C.S. */
Flag[P2] = false;
```

.....

## Mutual exclusion with busy waiting

**Problem:**

**deadlock: (Flag[P<sub>1</sub>]=true , Flag[P<sub>2</sub>]=true), infinite loop → rule 3 is broken!**

# Algorithm 3

- ▶ By combining the two algorithms, we obtain an algorithm which fulfils all three requirements for correctness
- ▶ In order to enter the critical section,  $P_i$  sets the `flag[ $P_i$ ]` to `true`, and sets `turn =  $P_j$` , which means that the other processes can enter their critical section (if they need to)
- ▶ The process that sets the flag and has its own value in the variable `turn`, enters the critical section



# Solution: Combination

Flag[P1..P2], turn

Process 1

```
Flag[P1] = true;
turn = P2;
while(Flag[P2] and turn == P2)
    ;
/* enter C.S. */
/* leave C.S. */
Flag[P1] = false;
.....
```

Process 2

```
Flag[P2] = true;
turn = P1;
while(Flag[P1] and turn == P1)
    ;
/* enter C.S. */
/* leave C.S. */
Flag[P2] = false;
.....
```

**Peterson's Solution**

(modern version of Dekker's algorithm)

# Peterson's Solution

```
#include <prototypes.h>
#define FALSE 0
#define TRUE 1
#define N 2 // number of processes
turn int turn, interested[N] // all values initially 0
void enter_region(int process)
{
    int other; // number of the other process
    other = 1 - process;
    interested[process] = TRUE; // show that you are
    interested
    turn = other; // set the flag
    while ( turn == other && interested[other] == true) ; // wait
}
void leave_region(int process) { // process who is leaving
    interested[process] = FALSE; // indicate departure from C.S.
}
```





# Peterson's Solution (Example)

Process = 0

```
interested[0]=interested[1]=FALSE
int turn;
int interested[2];

void enter_region(...){
    int other = 1;
    interested[0] = TRUE;
    turn = 1;
    while (turn == 1 &&
           interested[1] == TRUE);
}

void leave_region(...){
    interested[0] = FALSE;
}
```

Process = 1

```
interested[0]=interested[1]=FALSE
int turn;
int interested[2];

void enter_region(...){
    int other = 0;
    interested[1] = TRUE;
    turn = 0;
    while (turn == 0 &&
           interested[0] == TRUE);
}

void leave_region(...){
    interested[1] = FALSE;
}
```



## 2. Locking: Hardware support

- ▶ Disabling interrupts
- ▶ Special machine instructions



## 2. Disabling interrupts

- ▶ A simple solution for a single-processor systems:
  - Disable interrupts before the critical section
  - Enable interrupts after the critical section
    - When interrupts are disabled, no clock interrupt can occur -> there is no context switch -> there is no time limit for the process
- ▶ Disabling HW interrupts is a dangerous manoeuvre for the OS:
  - It can postpone the system response to important events
  - A process may never exit, leaving the interrupts blocked!
- ▶ Unwise solution for multi-processor systems
  - Disabling interrupts affects only one core, while other cores may access the shared variable
- ▶ It is not suitable as a technique for user processes
  - Convenient approach for the kernel to disable interrupts for a few instructions, while it is updating variables or lists
  - It does not work for systems with multiple CPUs (CPU cores)



# Example: Disabling Interrupts

shared double balance;

## Process P1

```
disableInterrupts();  
balance = balance + amount;  
enableInterrupts();
```

## Process P2

```
disableInterrupts();  
balance = balance - amount;  
enableInterrupts();
```

- ▶ The interrupts can be disabled an arbitrarily long time
- ▶ We just want  $P_1$  и  $P_2$  not to interfere; using this technique we block all  $P_i$
- ▶ We can use a mutual “lock” variable



## 2. Special Atomic Instructions

- ▶ Special atomic machine instructions
  - CPU executes them in a **single** instruction cycle
  - They are not subject to instruction mixing
  - They are used for control of access to a memory location:
    - Reading and writing
    - Reading and testing



# Test and Set Lock Instruction

- ▶ It executes **atomically**, in **one cycle**, **without interrupts**
  - TSL RX,LOCK reads the contents of the shared memory word **lock** into register **RX** and then stores a nonzero value at the memory address **lock**.
    - No other processor can access the memory word until the instruction is finished.
    - The CPU executing the instruction locks the memory bus to prohibit other CPUs from accessing memory until it is done.

enter\_region:

TSL REGISTER,LOCK	copy lock to register and set lock to 1
CMP REGISTER,#0	was lock zero?
JNE enter_region	if it was non zero, lock was set, so loop
RET	return to caller; critical region entered

leave\_region:

MOVE LOCK,#0	store a 0 in lock
RET	return to caller

# Test and Set Lock Instruction– Realization

- ▶ This instruction cannot be interrupted – no one can access the memory while it is executing

```
boolean TestSet (int i) {  
    if (i == 0) {  
        i = 1;  
        return true;  
    }  
    else {  
        return false;  
    }  
}
```



# Usage

- ▶ Shared boolean variable lock, initially set to false.
- ▶ Solution:

```
int lock=false;
while(true){
    while(!TestSet(&lock))
        ; /* do nothing
    //      critical section
    lock = FALSE;
    //      uncritical section

}
```

```
boolean TestSet(int i) {
    if (i == 0) {
        i = 1;
        return true;
    }
    else {
        return false;
    }
}
```



# Exchange of Values

- ▶ Exchange instruction (XCHG)
- ▶ An atomic exchange between a register and a memory location
- ▶ Atomic operation, in a single CPU instruction

enter\_region:

MOVE REGISTER,#1

XCHG REGISTER,LOCK

CMP REGISTER,#0

JNE enter\_region

RET

| put a 1 in the register

| swap the contents of the register and lock variable

| was lock zero?

| if it was non zero, lock was set, so loop

| return to caller; critical region entered

leave\_region:

MOVE LOCK,#0

RET

| store a 0 in lock

| return to caller



# XCHG Realization

```
void exchange(int register, int memory) {  
    int temp;  
    temp = memory;  
    memory = register;  
    register = temp;  
}
```



# Machine Instructions for Mutual Exclusion

## ▶ Pros

- Can be used for an arbitrary number of processors which can share the same working memory
- They are simple for work and verification
- Can be used for several critical sections



# Machine Instructions for Mutual Exclusion

## ► Cons

- Busy waiting wastes CPU time
- Starvation is possible when a process exits from the critical section, but more than one process are waiting to enter



# Basic Problem

- ▶ These approaches (with busy waiting) are wasting CPU time while waiting for a permission to enter the critical section
- ▶ Example (Priority Inversion Problem)
  - 2 processes with different scheduling priorities:
    - H – with high priority (the CPU is assigned to this process when it is in ready state)
    - L – with low priority
  - Scenario
    - L is executing the critical section
    - H becomes ready and the CPU is assigned to it (because it has higher priority).
    - H begins with busy waiting because L is in the critical section (locked the access)
    - L cannot leave the critical section because H “is working” (even though it waits), and therefore, is never scheduled
    - H is in an infinite loop




# sleep(), wakeup()

- ▶ Solution: Use of system calls which block the process instead of busy waiting
- ▶ Example:
  - `sleep()` suspends the process until some other process awakes him with `wakeup()`
  - `wakeup()` has one parameter, the process that needs to be awoken
- ▶ System calls of this type are used in solutions with blocking



# Bounded Buffer Problem

- ▶ Producer-consumer problem
- ▶ Important for: network interfaces, I/O controllers, message exchange, etc.



- ▶ Bounded buffer with limited capacity (N) shared by many processes
  - The producer puts information into the buffer (on the tail)
  - The consumer takes information out of the buffer (from the head)



# Bounded Buffer Problem

## ▶ Possible problems:

- The producer wants to put a new item in the buffer, but it is already full.
- The consumer wants to remove an item from the buffer, but the buffer is empty.

## ▶ Simple solution:

- Full buffer for the producer: it goes to sleep until the consumer removes one or more items and wakes it up.
- Empty buffer for the consumer: it goes to sleep until the producer puts something in the buffer and wakes it up.





# Solution: Classical Blocking

```
#include <prototypes.h>
#define N 100 // number of slots
int count 0;
void producer (void) {
    int item;
    while (TRUE) {
        produce_item(&item); //next item
        if (count == N) sleep(); //full?
        insert_item(item); //put item
        count = count + 1;
        if (count == 1)
            wakeup(consumer); // empty?
    } }

void consumer (void) {
    int item;
    while (TRUE) {
        if (count == 0) sleep(); //empty?
        remove_item(&item); // take item
        count = count - 1;
        if (count == N-1)
            wakeup(producer); // full?
        consume_item(item); // print item
    }
}
```

# Realization

- ▶ Problem: a race condition because access to *count* is unconstrained

Producer	Consumer
	Test of <code>count==0</code> , read <code>count</code> (just before <code>sleep()</code> )
<code>insert_item()</code>	
<code>count++</code>	
<code>wakeup(consumer)</code>	consumer is not sleeping (signal is lost)
produce next item	
	<code>count==0 =&gt; sleep</code>
<code>insert_item()</code>	
...	
<code>count==N =&gt; sleep</code>	



# Questions?

