

Deadlocks

Operating Systems

Assoc. Prof. Milos Jovanovik, PhD

Objectives

- ▶ Resources – system model
- ▶ Necessary conditions for deadlocks
- ▶ Detection
 - Resource-allocation graphs
- ▶ Handling methods
 - Prevention
 - Avoidance
 - Deadlock detection and recovery
 - Do nothing



System Model: System Resources

- ▶ The system consists of a finite number of resources, distributed among processes that compete to acquire them.
- ▶ There are several resource types, and each one has several number of identical instances.
 - Memory space, CPU cycles, files, I/O devices (e.g. printers) are different resource types;
 - If the system has 2 CPUs, then we say that the resource from type CPU has 2 instances. Similarly, the printer can have 5 instances;
- ▶ If the process requires a resource instance, then any instance should satisfy the requirement
 - If not, it cannot be regarded as the same resource;



Deadlock Problem

- ▶ A scenario in which 2 or more processes wait on each other endlessly.
- ▶ A set of processes is blocked when one of them is waiting on an event which can be triggered only by another process from the same set.
- ▶ Can happen when using mutual exclusion (working in critical sections).
- ▶ Example:
 - Process A has the **printer**, but now requests for a **file**;
 - Process B has the **file**, but now requests for the **printer**.



An Example with Databases

Process A

`lock(r2);`

...


`lock(r1); // tries`

Process B

`lock(r1);`

...

`lock(r2); // tries`



Deadlocks are Unwanted

Because:

- ▶ They “stop” the process from progressing
- ▶ They require intervention
- ▶ They reduce resource utilization
- ▶ All processes are waiting:
 - No process will be able to trigger an event that will awake some other process;
 - Usually the process waits on some resource to be released, which is acquired by some other blocked process;
- ▶ NONE of the processes:
 - Works
 - Can release a resource
 - Can be awoken
- ▶ All processes are waiting infinitely



Conditions for a Deadlock

- ▶ **Mutual Exclusion:** Only one process at a time can use a resource.
- ▶ **Hold and Wait:** A process holding at least one resource, is waiting to acquire additional resources held by other processes.
- ▶ **No Preemption:** A resource can be released only voluntarily by the process holding it, after that process has completed its task.
- ▶ **Circular Wait:** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource held by P_1 , P_1 is waiting for a resource held by P_2 , ..., P_{n-1} is waiting for a resource held by P_n , and P_n is waiting for a resource that is held by P_0 .

All conditions must be met simultaneously!



Solutions

- ▶ In general, there are 4 strategies to solve deadlocks:
 - **Prevention**: the concurrency in the OS is implemented in such a way that a deadlock is impossible
 - **Avoidance**: foresee the deadlock and avoid it
 - **Allow** a deadlock and react
 - **Do nothing** (mostly used)



Systems with multiple Processes and Resources

- ▶ In order to detect a deadlock in the system, we need a **formalization** which will allow us to represent the processes and the resources in the system
- ▶ The symplest way is to represent the system as a **resource-allocation graph**




Resource–Allocation Graph

- ▶ A set of vertices V and a set of edges E
- ▶ V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- ▶ Request edge: directed edge $P_i \rightarrow R_j$
- ▶ Assignment edge: directed edge $R_j \rightarrow P_i$




Resource-Allocation Graph


- ▶ Process




- ▶ Resource with 4 instances



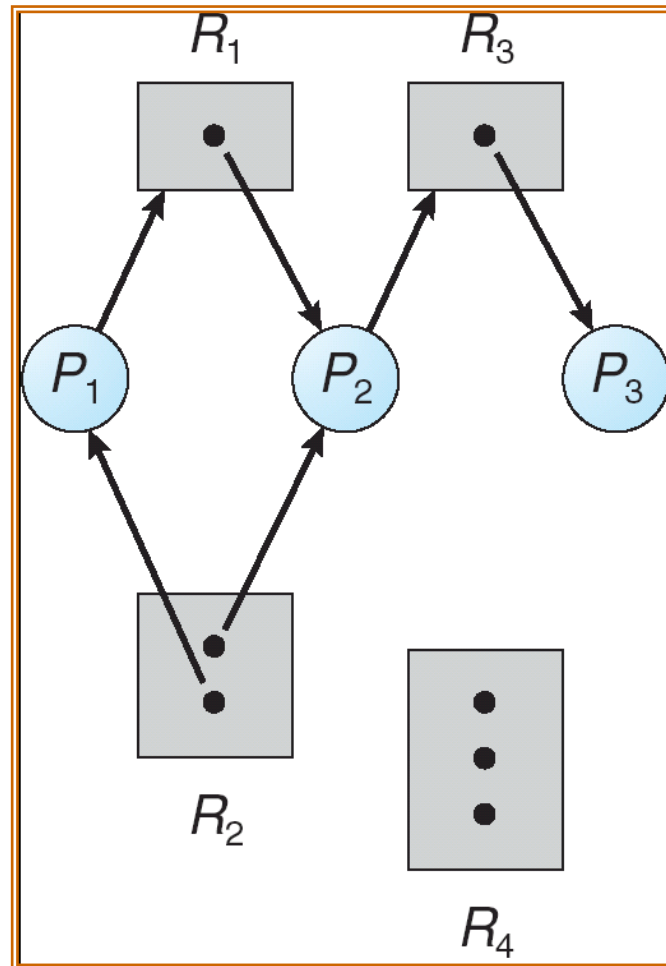
- ▶ P_i asks for one instance of R_j



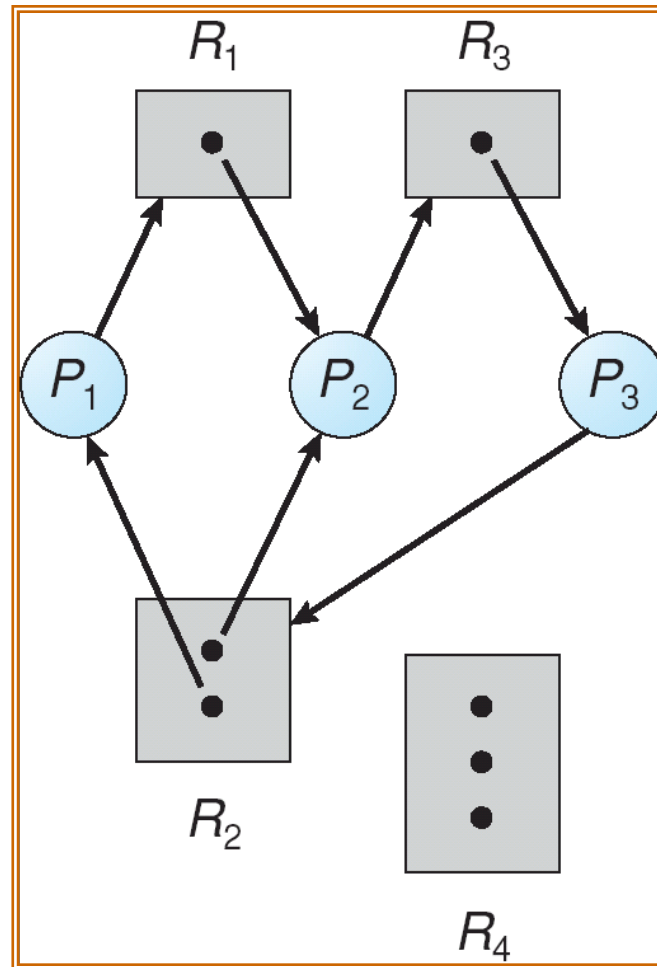
- ▶ P_i holds one instance of R_j




Resource-Allocation Graph: Example



Resource-Allocation Graph: Example with a Deadlock



Graph with a Cycle, but without a Deadlock



Basic Facts

- ▶ If graph contains no cycles \Rightarrow no deadlock.
- ▶ If graph contains a cycle \Rightarrow
 - If there is only one instance per resource type, then there is a deadlock.
 - If there are several instances per resource type, there is a possibility of deadlock.




Example

- ▶ Assume that you have the following system:
 - There are three types (classes) of resources: R1, R2, R3.
 - There are two instances of each resource type.
 - There are three processes: P1, P2, P3.
- ▶ Some of the resource instances have already been assigned to processes:
 - Two instances of R1 are allocated to P1;
 - One instance of R2 is allocated to P2, and the other one to P3;
 - One instance of R3 is assigned to P1, and the other one to P3;
- ▶ Some processes have requested new instances:
 - P3 has requested one instance of R1;
 - P2 has requested one instance of R1, as well;




Example

- ▶ Which process can finish first?




Example

- ▶ Which process can finish first?
 - P1 can finish first, it is in the “ready” state




Example

- ▶ What about P2 and P3?
 - They are in the state of “**blocked**”




Example

- ▶ What happens after P1 finishes?
 - There are two R1 instances for P2 and P3, each



Example

- ▶ So, is the system in a **deadlock**?
 - No, because all processes can finish



Prevention

How to prevent Deadlocks

Conditions for Mutual Exclusion

- ▶ If we deny exclusive resource assignment to some process --> NO deadlocks...
- ▶ Not required for sharable resources; However, we must apply mutual exclusion for non-sharable resources.
- ▶ Some resources are intrinsically non-sharable.
- ▶ This is used in OS cores.



Conditions for Hold and Wait

- ▶ There must be a guarantee that whenever a process requests a resource, it does not hold any other resources.
 - Require process to request and get all resources before they begin their execution, or allow a process to request resources only when the process has none.
 - Drawbacks:
 - Low resource utilization;
 - Starvation is possible;



Conditions for Preemption

- ▶ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- ▶ Preempted resources are added to the list of resources for which the process is waiting.
- ▶ The process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.



Conditions for Preemption

- ▶ This is the worst condition for preventing deadlocks:
 - If a process has a printer assigned, and is in the middle of the job, taking away the printer just because the needed plotter is currently unavailable is a bad (if not an impossible) solution;




Conditions for Circular Wait

- ▶ Rule: Impose a total ordering of all resources.
- ▶ Require that each process requests resources in an increasing order of enumeration.
 - A process that holds a resource, cannot ask for a resource with a smaller number, until it frees its own resource
- ▶ Keep in mind that developing an ordering, or hierarchy, does not in itself prevent deadlock.
 - It is up to application developers to write programs that follow the ordering;
- ▶ Problems:
 - Inflexible solution;
 - Not applicable for bigger systems;



Deadlock Avoidance

How to avoid Deadlocks



Banker's Algorithm

- ▶ Resources with multiple instances.
- ▶ Each process must a priori claim maximum use.
- ▶ When a process requests a resource, it may have to wait.
- ▶ When a process gets all its resources, it must return them in a finite amount of time.



Data Structures for the Banker's Algorithm

n = number of processes, and m = number of resources types.

- ▶ **Available**: Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- ▶ **Max**: $n \times m$ matrix. If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j .
- ▶ **Allocation**: $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j .
- ▶ **Need**: $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$



Resource-Request Algorithm for Process P_i

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:
 - $Available = Available - Request_i;$
 - $Allocation_i = Allocation_i + Request_i;$
 - $Need_i = Need_i - Request_i;$
 - If the new state is safe \Rightarrow the resources are allocated to P_i .
 - If the new state is unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored



Example of Banker's Algorithm

- ▶ 5 processes P_0 through P_4 ;
- ▶ 3 resource types
 - A (10 instances),
 - B (5 instances), and
 - C (7 instances).
- ▶ Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			



Example of Banker's Algorithm

- ▶ The content of *Need* is *Max - Allocation*.

	<u>Need</u>			<u>Available</u>		
	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>
P_0	7	4	3	3	3	2
P_1	1	2	2			
P_2	6	0	0			
P_3	0	1	1			
P_4	4	3	1			

- ▶ Safe sequences:
 - $\langle P_1, P_3, P_4, P_2, P_0 \rangle$
 - $\langle P_1, P_4, P_3, P_0, P_2 \rangle$
 - ...



Example Resource-Request Algorithm

- ▶ P1 Request (1,0,2)
- ▶ Check that Request \leq Available
 - $(1,0,2) \leq (3,3,2) \Rightarrow \text{true.}$

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	7	5	3	3	3	2
P ₁	2	0	0	3	2	2			
P ₂	3	0	2	9	0	2			
P ₃	2	1	1	2	2	2			
P ₄	0	0	2	4	3	3			



Example Resource-Request Algorithm

- ▶ State after allocating the requested resources:

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	7	4	3	2	3	0
P ₁	3	0	2	0	2	0			
P ₂	3	0	2	6	0	0			
P ₃	2	1	1	0	1	1			
P ₄	0	0	2	4	3	1			

- ▶ Banker's algorithm shows that sequence <P1, P3, P4, P0, P2> satisfies safety requirement.
 - Can request for (3,3,0) by P4 be granted? – not available
 - Can request for (0,2,0) by P0 be granted? – not safe




Usage of Deadlock Avoidance?

- ▶ Useful theoretically, but useless in practice
 - How will the system know the number of resources per process *in advance*?
 - The number of active processes changes dynamically, new users come and go!
 - Some resource can also be broken (physical error, driver error, etc.)!
- ▶ In practice, there is a very small number of systems that use the banker algorithm for deadlock avoidance!



Allow, Detect and React

React on a Deadlock



Allow, Detect and React

- ▶ Allow system to enter a deadlock state
- ▶ Analyse the situation (detection algorithm)
- ▶ Recover from deadlock
 - Pick a victim-process and kill it
- ▶ Questions:
 - How to detect when a deadlock is happening?
 - We need a formal algorithm: detecting cycles in oriented graph or a matrix algorithm;
- ▶ How to pick the victim?
 - To recognize the right process that makes the deadlock...

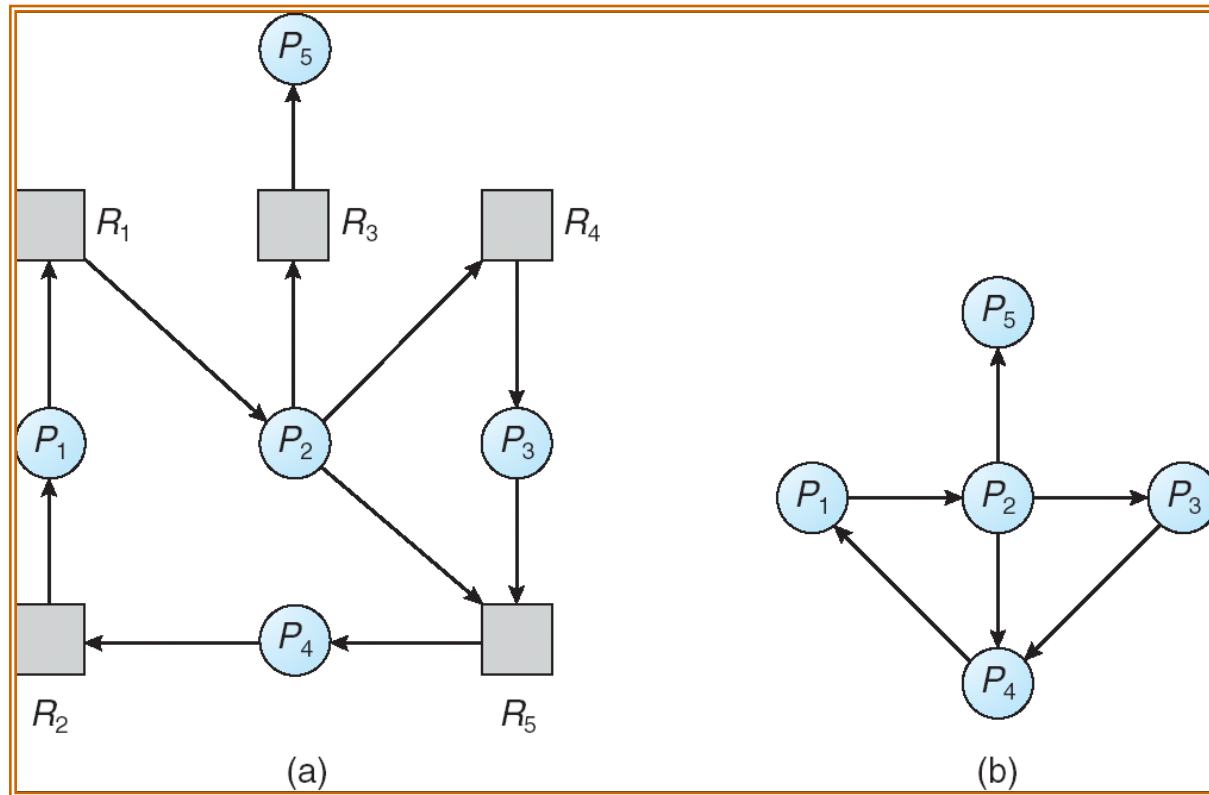


Single Instance of Each Resource Type

- ▶ Maintain a *wait-for* graph
 - Nodes are processes.
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j .
- ▶ Periodically invoke an algorithm that searches for a **cycle in the graph**. If a cycle is there, there's a deadlock.
- ▶ An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph.



Resource-Allocation Graph and Wait-For Graph



Resource-Allocation Graph

Corresponding Wait-For Graph

Do Nothing!

Don't prevent, avoid or
proactively detect Deadlocks

Do not react!

- ▶ The simplest algorithm
- ▶ Ostrich algorithm (head in the sand)
 - How often there is a deadlock?
 - Mathematician vs Engineer!



Balance: Correctness vs. Comfort

- ▶ In each OS, the tables are finite and that leads to problems
 - Process and thread problems
 - Open files
 - Swap space
- ▶ It is probable, that each user prefers rare deadlocks over restrictions on resource use
 - **Balance between correctness and commodity**



Windows 11: Deadlocks

- ▶ **Resource Allocation Graph:** Windows 11 uses the RAG algorithm to detect deadlocks.
 - This algorithm represents the resource allocation relationship among processes as a directed graph, and checks for cycles in the graph;
 - If a cycle exists, it indicates a deadlock;
- ▶ **Timeout Mechanism:** Windows 11 employs a timeout mechanism to prevent deadlocks from occurring.
 - If a process is unable to acquire a resource within a specified time period, it releases all the resources it has acquired and goes into a wait state;
 - This ensures that a process does not hold resources indefinitely;
 - The timeout mechanism in Windows is implemented by setting a maximum wait time for processes to acquire resources and checking for cycles in the wait-for graph to detect deadlocks;



Windows 11: Deadlocks

- ▶ **Preemption:** Windows 11 uses preemption to resolve deadlocks.
 - If a process holds a resource and is unable to acquire another resource, the OS may preempt the resource and allocate it to another process;
 - This mechanism ensures that resources are used efficiently and prevents deadlocks from occurring;
- ▶ **Deadlock Detection and Recovery:** Windows 11 has built-in mechanisms to detect and recover from deadlocks.
 - When a deadlock is detected, the OS may terminate one or more processes involved in the deadlock to break the cycle and recover from the deadlock;
 - Random Process Termination, Priority-Based Process Termination, Resource-Usage-Based Process Termination, Deadlock Avoidance (is implemented to prevent the deadlocks – BY NOT GRANTING THE RESOURCE);



Linux: Deadlocks

- ▶ **Lock ordering:** The Linux kernel uses a strict lock ordering protocol to prevent deadlocks.
 - Lock ordering means that locks are always acquired in a specific order, and released in the reverse order;
 - This ensures that no process can acquire a lock that is held by another process, preventing circular dependencies and deadlocks;
- ▶ **Lockdep:** The Linux kernel includes a tool called "lockdep" that checks for potential deadlocks at runtime.
 - Lockdep analyses the lock dependencies in the kernel code and reports any potential circular dependencies;
 - This helps developers to identify and fix potential deadlocks before they occur;



Linux: Deadlocks

- ▶ **RCU (Read-Copy-Update):** RCU is a synchronization mechanism used in the Linux kernel to allow concurrent read access to shared data structures without the need for locks.
 - RCU uses a deferred deletion mechanism to remove data structures, ensuring that no process can access a deleted data structure;
- ▶ **Wait-for graph analysis:** The Linux kernel uses a wait-for graph to detect deadlocks.
 - The wait-for graph is a directed graph that represents the dependencies between processes and the resources they require;
 - When a process is waiting for a resource, it is added as a node to the wait-for graph;
 - If a process is waiting for another process that is also waiting for a resource, a cycle is detected, and a deadlock is assumed;
 - In this case, the kernel takes action to resolve the deadlock by either releasing the resources or terminating one of the processes involved;



Linux: Deadlocks

- ▶ **Priority inheritance:** Priority inheritance is a mechanism used in the Linux kernel to prevent priority inversion, which can lead to deadlocks.
 - Priority inheritance ensures that the priority of a low-priority process is temporarily increased to the priority of the highest-priority process that is waiting for a resource held by the low-priority process;
 - This prevents the high-priority process from being blocked and potentially causing a deadlock;



Conclusion

- ▶ Deadlocks are a **potential problem** for every system
- ▶ They can be **avoided** with **safe states** (there is a list of events that guarantees that all the processes will finish)
- ▶ The Banker algorithm **avoids** deadlocks by not allowing the request to force the system in an unsafe state
- ▶ The deadlock can be avoided with enumerating, so that each processes asks for resources in increasing order
- ▶ The starvation can be avoided with FCFS policy



Questions?

