


Data Storage

Operating Systems

Assoc. Prof. Milos Jovanovik, PhD

Stable Storage



Stable Storage

- ▶ When a write is issued to the disk, it should either correctly write the data, or do nothing (leaving the existing data intact)
- ▶ Reasons:
 - ECC is not enough
 - A correctly written sector can spontaneously go bad and become unreadable
 - CPU can fail
- ▶ A pair of identical disks is used
- ▶ Implemented in software



Stable Storage

▶ Stable write

- The data is written and verified (read) on both disks (first HD1, then HD2)
- If the data is not verified, the operation is repeated up to n times
- After n consecutive failures, the block is remapped onto a spare and the operation gets repeated until it succeeds

▶ Stable read


- Reads the block from HD1 and checks ECC
- If this yields an incorrect ECC, the read is tried again, up to n times.
- If all of these give bad ECCs, the corresponding block is read from HD2

▶ Crash recovery

- If a pair of blocks are both correct, nothing is done
- If one of them has an ECC error, the bad block is overwritten with the corresponding good block
- If a pair of blocks are both good but different, the block from HD1 is written onto HD2



Stable Storage




Analysis of the influence of crashes on stable writes.

RAID

- ▶ In order to satisfy the demand for higher rates and parallel data processing, several different architectures are used for disk organization – RAID (Redundant Array of Inexpensive/Independent Disks).
- ▶ The basic concept is to connect several disks in order to achieve redundancy and availability.
- ▶ Transparent to the processor.




RAID 0 – Strip Set



- ▶ Each strip contains k sectors.
- ▶ It is used to increase the performances – parallel I/O.
- ▶ The capacity depends on the smallest disk:
 - HD1 = 100GB, HD2 = 120 GB, Total = 200GB
- ▶ Negative parameter: $MTTF = MTTF / \#HD$.




RAID 1



- ▶ It duplicates all the disks, so there are four primary disks and four backup disks.
- ▶ On a write, every strip is written twice (bad performance).
- ▶ On a read, either copy can be used, distributing the load over more drives (good performance).
- ▶ Fault tolerance is excellent.




RAID 2



- ▶ Data is divided on a bit level and Hamming code is used for error correction.
- ▶ Possible high rates, but not used that often.
- ▶ We would need 39 disks (32 disks for a word, 7 for Hamming code and parity).
- ▶ This scheme requires all drives to be rotationally synchronized.




RAID 3



- ▶ RAID 3 consists of byte-level striping with dedicated parity.
- ▶ Although implementations exist, RAID 3 is not commonly used in practice.
- ▶ Error correction is possible
- ▶ The number of separate I/O requests per second they can handle is no better than for a single drive




RAID 4



- ▶ Strip-for-strip parity written onto an extra drive.
- ▶ It can serve several requests.
- ▶ The parity strip is obtained from the other 4 stripes with an XOR operation.
- ▶ If a drive crashes, the lost bytes can be recomputed from the parity drive by reading the entire set of drives.
- ▶ Negative side: When writing on one sector, all the disks needs to be read in order to update the corresponding parity stripe.



RAID 5

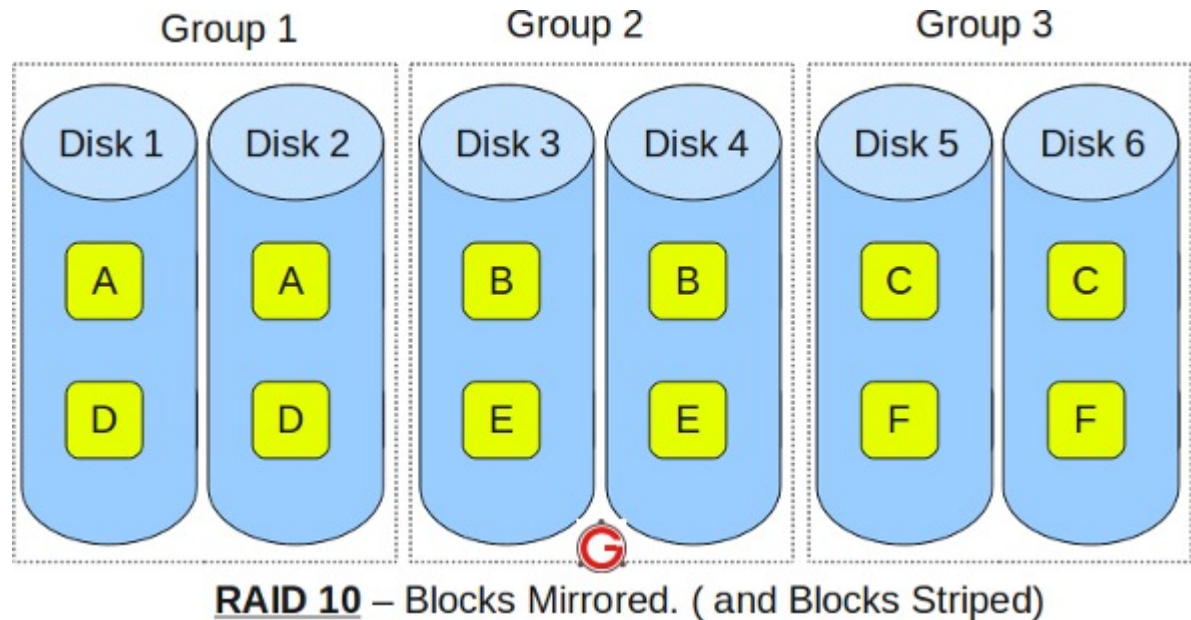


- ▶ As a consequence of the heavy load on the parity drive, it may become a bottleneck.
- ▶ This bottleneck is eliminated in RAID level 5 by distributing the parity bits uniformly over all the drives, in a round-robin fashion.
- ▶ However, in the event of a drive crash, reconstructing the contents of the failed drive is a complex process.



RAID 10

RAID 10 == RAID 1+0




- ▶ It needs at least 4 disks
- ▶ The disks are grouped in pairs as mirrored disks
- ▶ For 6 disks in RAID 10, there are 3 groups
- ▶ In one group, the data are mirrored – Disk 1 and Disk 2 are in group 1
- ▶ In the group, the data is stretched, i.e. Block A is written in group 1, Block B in Group 2, Block C in group 3.
- ▶ It is called “stripe of mirrors”, i.e. the disks in the group are mirrored and the data spans across groups.



Solid State Drives (SSDs)

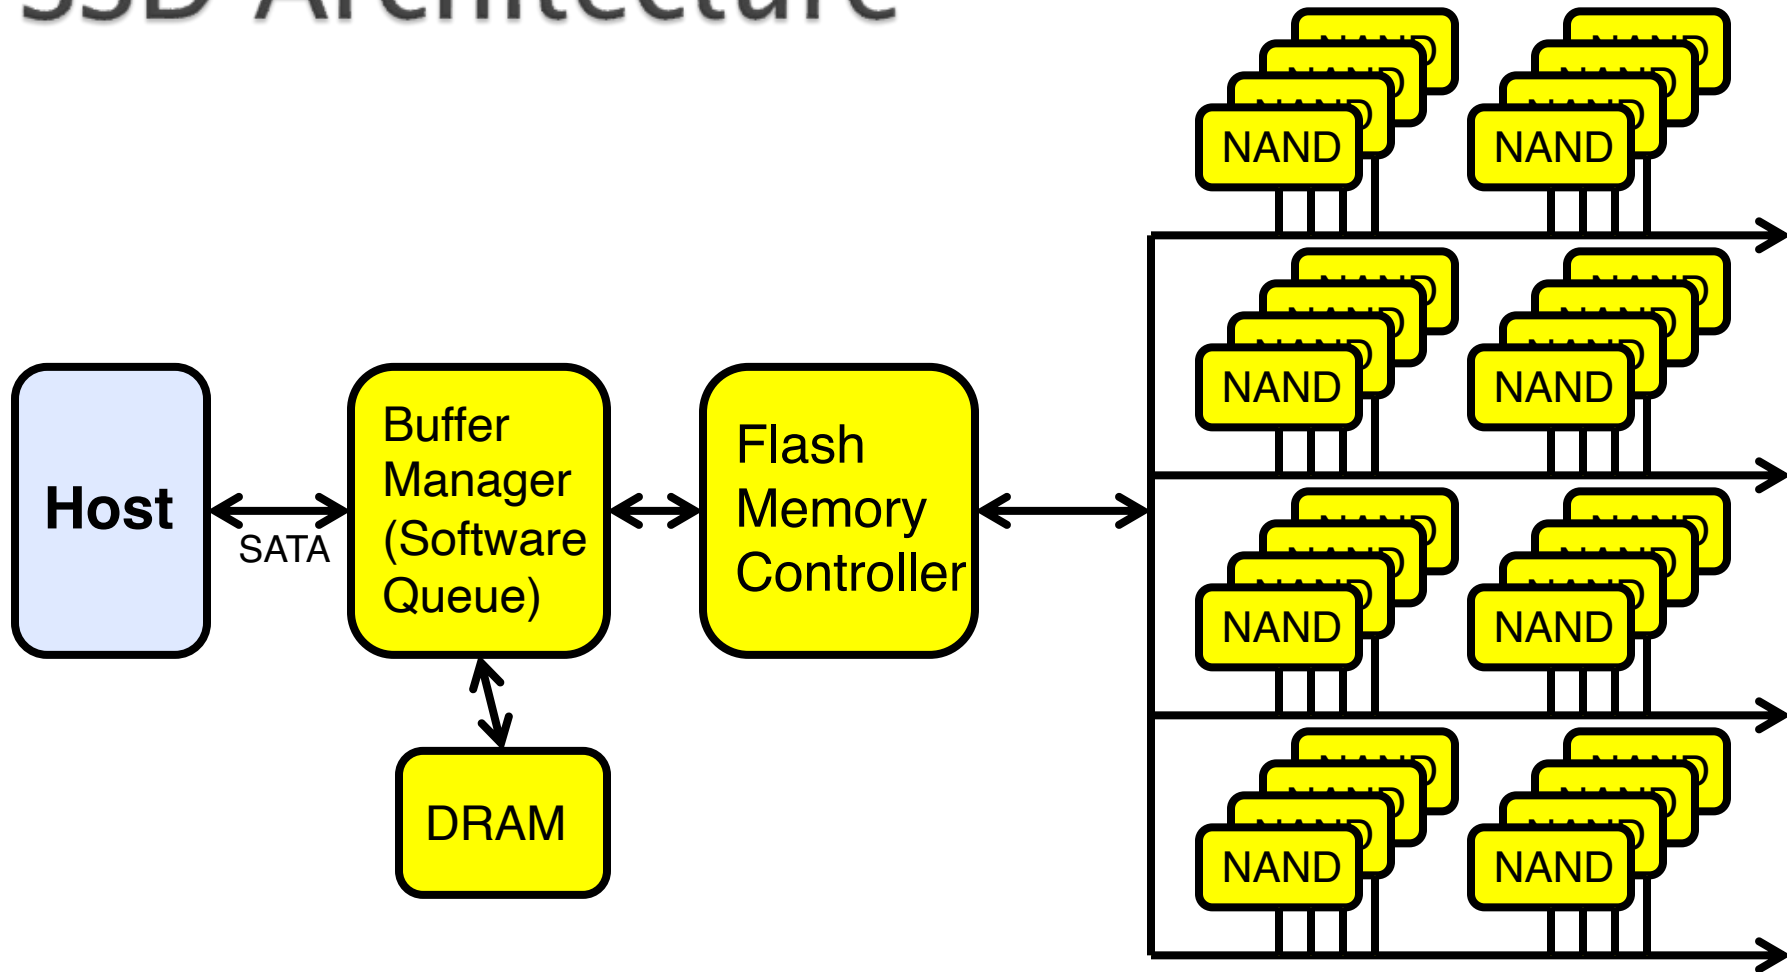
Solid State Disks (SSDs)



- ▶ 1995 – Replace rotating magnetic media with non-volatile memory (battery backed DRAM)
 - 2009 – Use NAND Flash: Single-Level Cell (1-bit/cell), Multi-Level Cell (2-bit/cell)
- ▶ Sector (4 KB page) addressable, but stores 4–64 “pages” per memory block
- ▶ No moving parts (no rotate / seek motors)
 - Eliminates seek and rotational delay (0.1–0.2ms access time)
 - Very low-power and lightweight



SSD Architecture

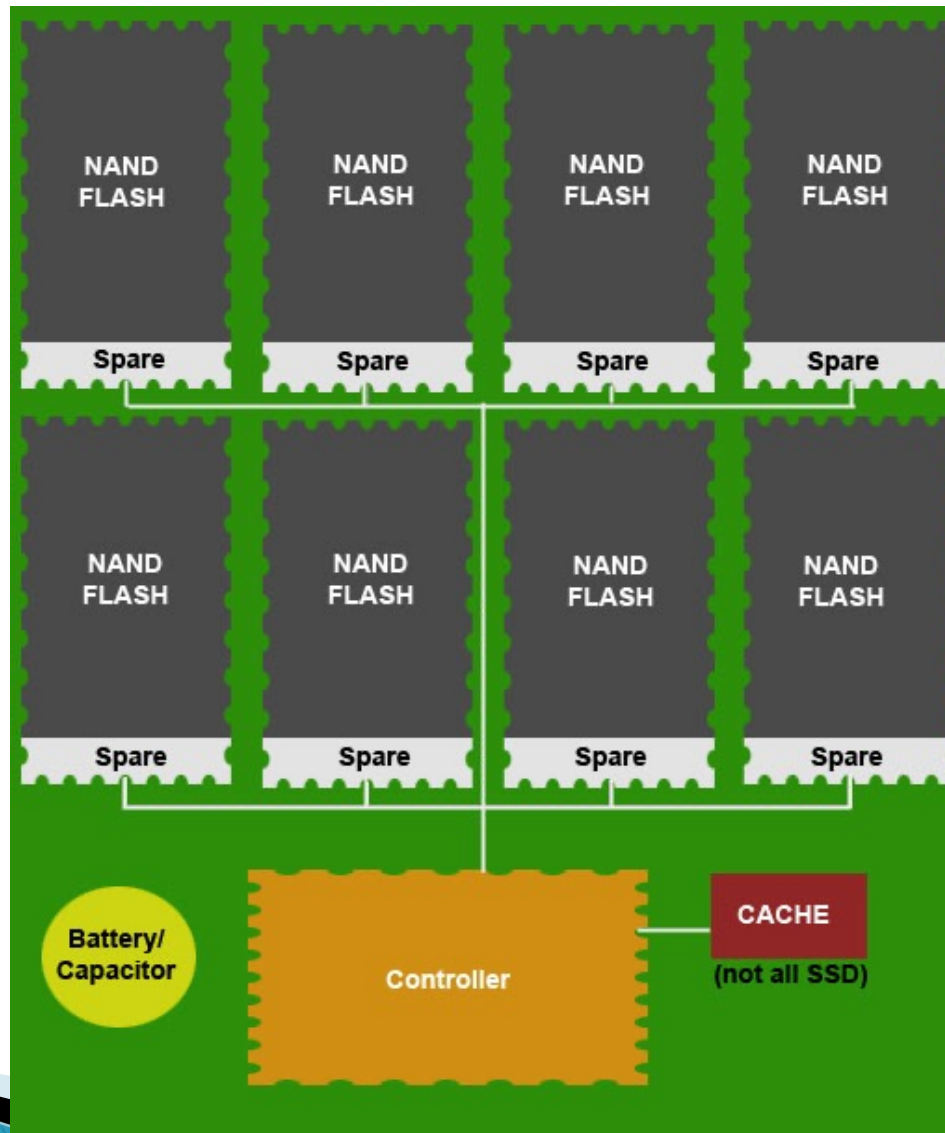


SSD Architecture

- ▶ SSD is divided into **blocks**, in which **pages** are kept.
- ▶ Pages consist of adjacent cells from NAND Flash memory.
- ▶ Blocks are grouping pages, and the number of blocks limits the size of the SSD.
- ▶ Pages are 4KB and blocks have 64 pages.
- ▶ **Data is written page-by-page, but deleted only block-by-block.**
 - Data is erased from 256KB ($64 \times 4\text{KB}$) blocks.

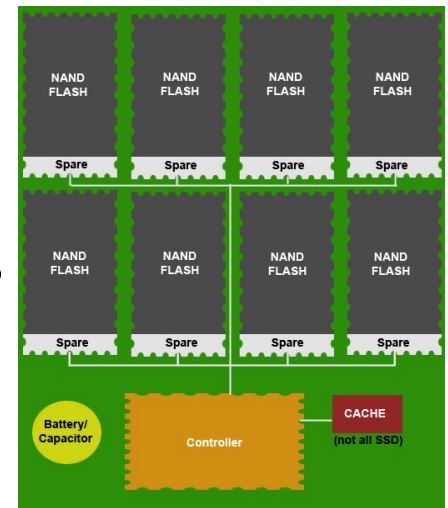


SSD Architecture Example



SSD Architecture Example

- ▶ 128 GB SSD with 8 NAND Flash chips
- ▶ Each chip is 16 GB or total of 120 GB, because 1GB is a spare chip
- ▶ The cache is not always present, and it has directories of which blocks are occupied and wear leveling
- ▶ Usually there are 4 to 10 channels to the NAND chips, which increases the throughput.
- ▶ Battery / Capacitor is used when there are problems with the power supply, in order to finish the process of writing on the SSD




SSD: Read

- ▶ Reading data similar to memory read (25μs)
- ▶ No seek or rotational latency
- ▶ Transfer time: transfer a 4KB page
 - Limited by controller and disk interface (SATA: 300–600MB/s)
- ▶ Latency = Queuing Time + Controller time + Transfer Time
- ▶ Highest Bandwidth: Sequential OR Random reads



Writing on SSD

- ▶ For SLC (single level cell), the cell is either On or Off.




- Writing is easier, there is less work for the controller




Writing on SSD

- ▶ Multi-level cell, the cell is On On, On Off, Off On or Off Off.
- ▶ It doubles the capacity
- ▶ It compromises reliability




Writing on SSD

- ▶ For instance, if the block size is 12 pages, each with size of 1B



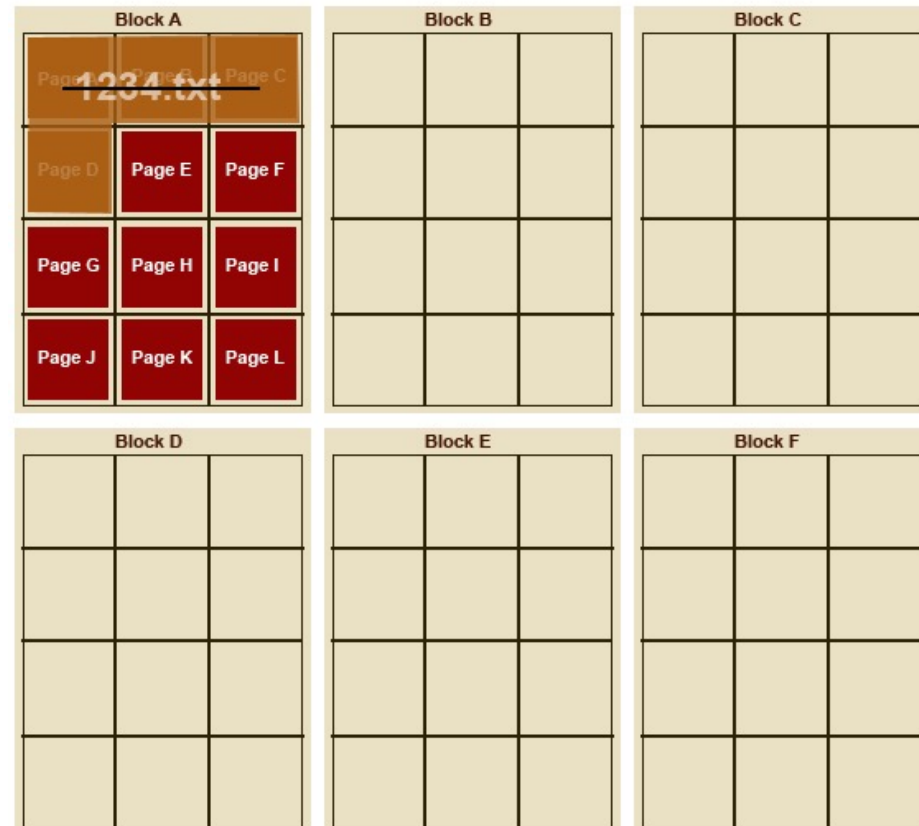
Writing on SSD

- ▶ Let us assume a file of 4B -> 4 pages



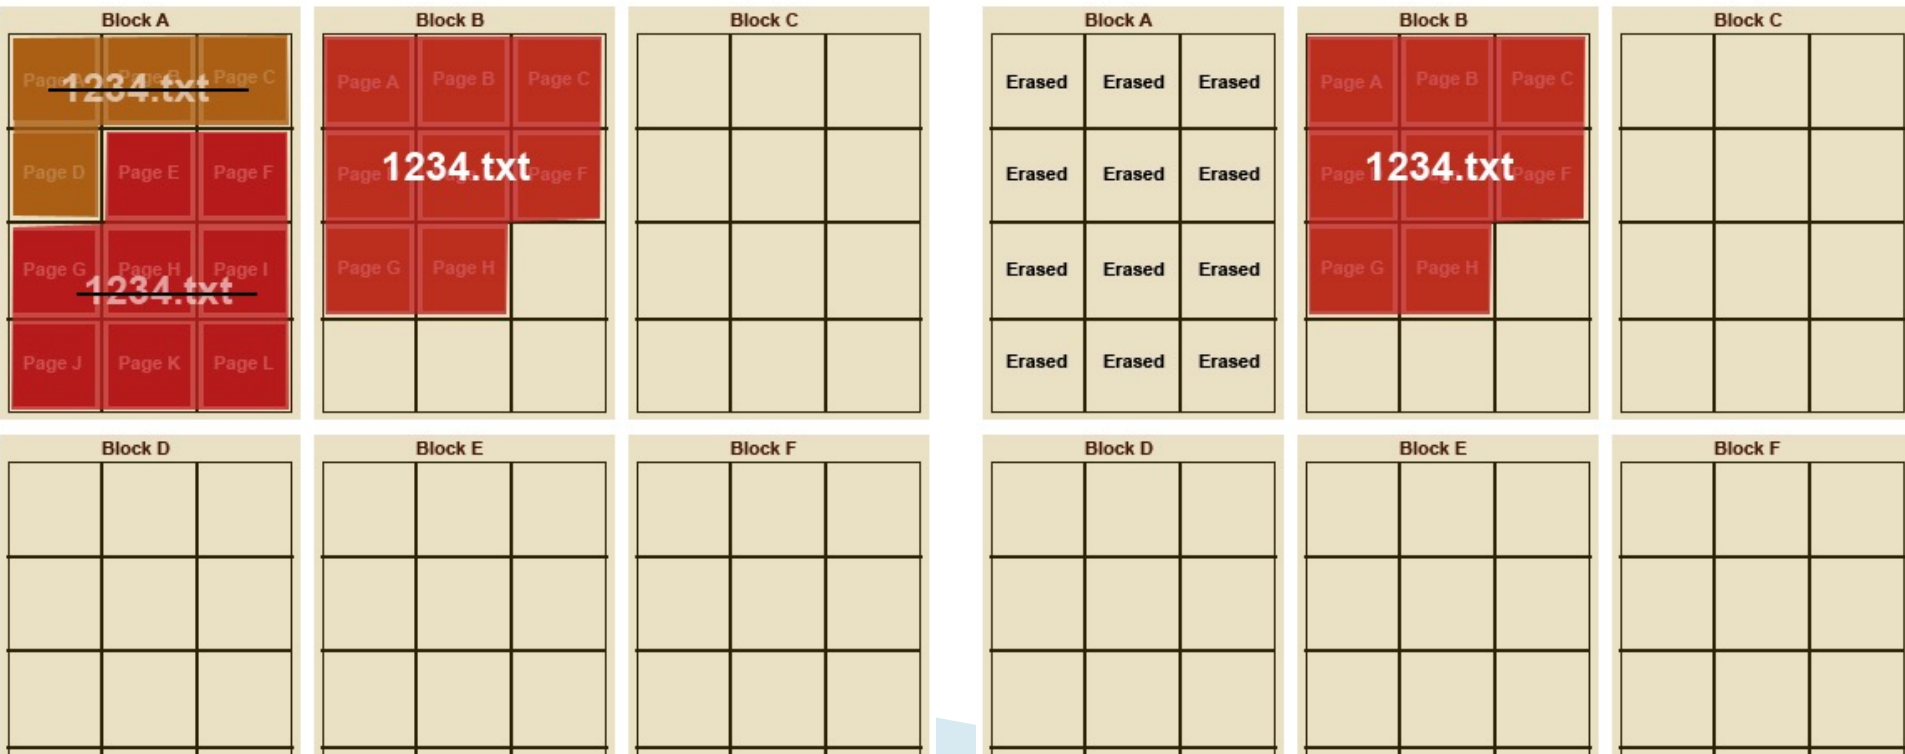
Writing on SSD

- ▶ If we change the file and we add text with size of 8B.
- ▶ Old pages are marked for deleting, whereas the new pages are written in the block



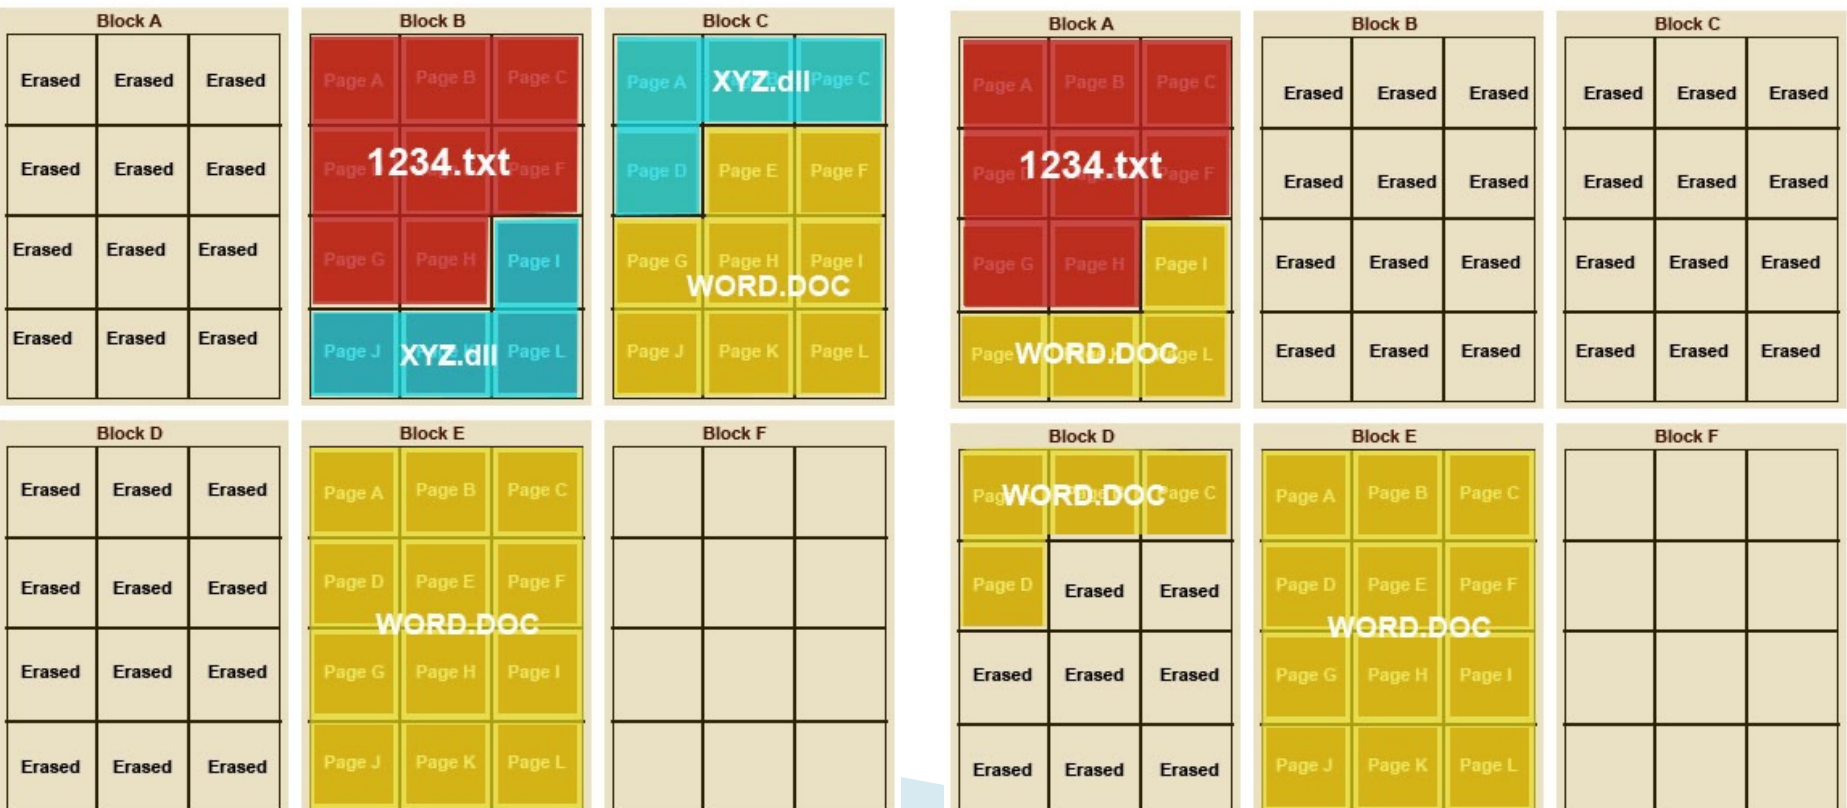
Writing on SSD

- ▶ SSD erases only full blocks, not pages!
- ▶ In order to erase old data, the new data must be written on block B, in order to erase data from block A



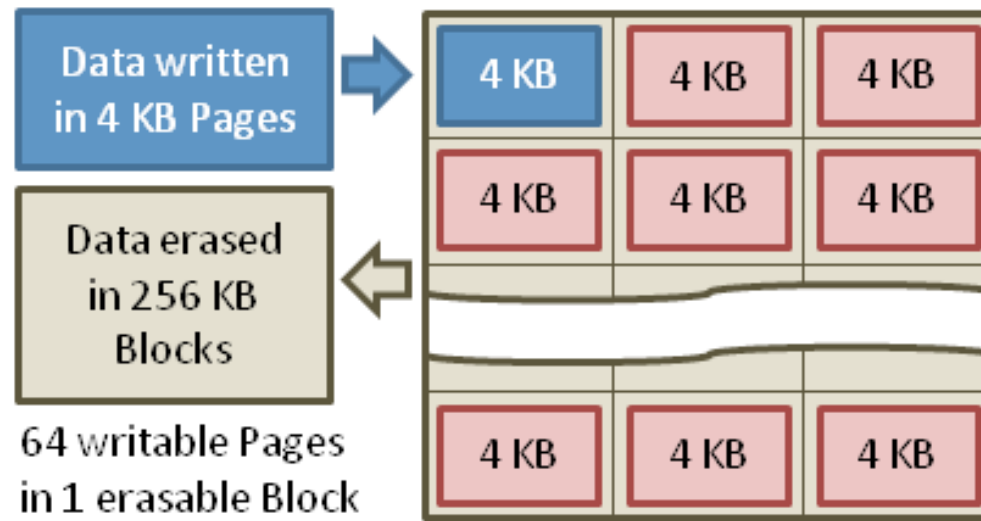
Writing on SSD

- ▶ What happens if we delete the file xyz.dll?



Writing on SSD

- Writing data is complex! ($\sim 200\mu\text{s}$ – 1.7ms)
 - Only empty pages can be written in a block
 - Only entire block can be erased ($\sim 1.5\text{ms}$)
 - Controller maintains pool of empty blocks by coalescing used pages (read, erase, write)
 - reserves some % of capacity

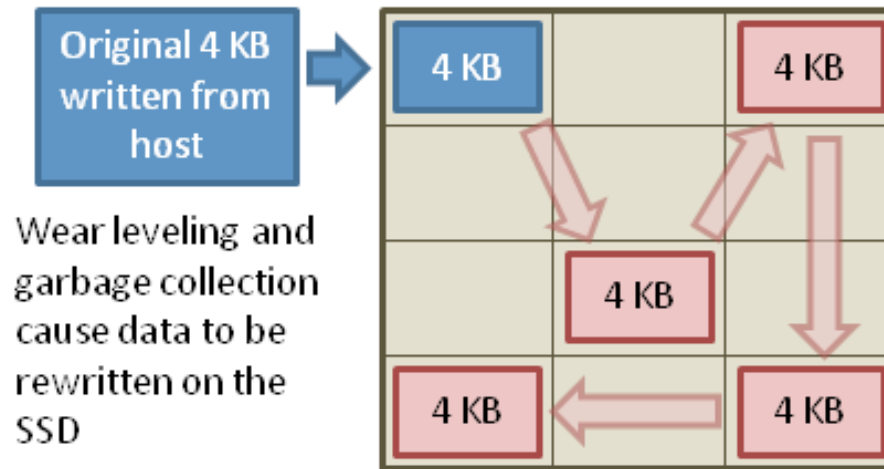


Typical NAND Flash Pages and Blocks



Writing on SSD

- ▶ Write and erase cycles require “high” voltage
 - Damages memory cells, limits SSD lifespan
 - Controller uses ECC, performs wear leveling



- ▶ Result is very workload-dependent performance
 - **Latency** = Queuing Time + Controller time (Find Free Block) + Transfer Time
 - **Highest BW**: Seq. OR Random writes (limited by empty pages)

Rule of thumb: writes 10x more expensive than reads, and erases 10x more expensive than writes

Writing on SSD

- ▶ Data cleaning is called **Garbage Collection** and keeps the SSD performances
- ▶ You never defragment SSD!
- ▶ The process of cleaning the SSD, before the data is written can be a downside, because it adds write activities which can slow the SSD's performance.
- ▶ You can write limited times on SSD




Writing on SSD

- ▶ Write A, B, C, D

| | | | |
|---------|------|------|------|
| Block X | A | B | C |
| | D | free | free |
| | free | free | free |
| | free | free | free |

Writing on SSD

- ▶ Write A, B, C, D
- ▶ Write E, F, G, H and A', B', C', D'
 - Record A, B, C, D as obsolete



Writing on SSD

- ▶ Controller *garbage collects* obsolete pages by copying valid pages to new (erased) block
- ▶ Typical steady state behavior when SSD is almost full
 - One erase every 64 or 128 writes



Storage Performance & Price

| | Bandwidth (Sequential R/W) | Cost/GB | Size |
|-----------------------|--|----------------|------------|
| HDD ² | 50-100 MB/s | \$0.03-0.07/GB | 2-4 TB |
| SSD ^{1,2} | 200-550 MB/s (SATA) 6 GB/s (read PCI) 4.4 GB/s (write PCI) | \$0.87-1.13/GB | 200GB-1TB |
| DRAM ² | 10-16 GB/s | \$4-14*/GB | 64GB-256GB |
| *SK Hynix 9/4/13 fire | | | |

¹<http://www.fastestssd.com/featured/ssd-rankings-the-fastest-solid-state-drives/>

²<http://www.extremetech.com/computing/164677-storage-pricewatch-hard-drive-and-ssd-prices-drop-making-for-a-good-time-to-buy>

BW: SSD up to x10 than HDD, DRAM > x10 than SSD

Price: HDD x20 less than SSD, SSD x5 less than DRAM



SSD Summary

- ▶ Pros (vs. HDD):
 - Low latency, high throughput (eliminate seek/rotational delay)
 - No moving parts:
 - Very lightweight, low-power, silent, very shock insensitive
 - Read at memory speeds (limited by controller and I/O bus)



SSD Summary

► Cons

- Small storage (0.1–0.5x HDD), very expensive (20x HDD)
 - Hybrid alternative: combine small SSD with large HDD
- Asymmetric block write performance: read /erase/write
 - Controller garbage collection (GC) algorithms have major effect on performance
- Limited drive lifetime
 - 1–10K writes/page for MLC NAND
 - Avg failure rate is 6 years, life expectancy is 9–11 years



Questions?

